

Distributed programming in Haskell

Mathieu Boespflug



Tweag I/O


Getting started...

```
$ git clone https://github.com/mboes/zurihac15-talk-examples  
$ cd zurihac15-talk-examples  
$ cabal sandbox init  
$ cabal update  
$ cabal install
```

Plan

1. Introduction
2. Cloud Haskell from the ground up
3. Code along session
4. Discussion





Distributed programming

- ★ From loosely coupled SOA...
 - REST API, closed set of verbs/resources
 - HTTP request/response
 - Extensible and/or implicit schema (JSON, XML, ...)
- ★ ...to tightly coupled distributed systems...
 - RPC, open set of verbs
 - message passing
 - fixed, explicit schemas (strong client/server coupling)
 - ... **Focus of this talk**

Why use Haskell?

- ★ (1) Strong library support (warp, http-client, aeson...)
- ★ (2) Strong library support (Cloud Haskell, zmq, ...)
- ★ The performance of a mature native compiler
- ★ Typeful programming for robust systems
- ★ Exceptionally powerful array of composition mechanisms for boilerplate-less, high-level programming.

Moving up from the nuts and bolts

★ Map/Reduce:

- map function

$$\text{MR.map} :: (k_{\text{in}}, v_{\text{in}}) \rightarrow [(k_{\text{med}}, v_{\text{med}})]$$

- reduce function

$$\text{MR.reduce} :: (k_{\text{inter}}, [v_{\text{inter}}]) \rightarrow [(k_{\text{out}}, v_{\text{out}})]$$

- **Reality:**

- long map/reduce pipelines (Google indexing pipeline: 21 step)
- 21 steps means 21 mapper classes and 21 reducer classes in Java

★ Two directions for improvement:

- more general data-parallel compute model
- less boilerplate!

WordCount Spark (Java)

```
JavaRDD<String> textFile = spark.textFile("hdfs://...");
```

```
JavaRDD<String> words = textFile.flatMap(new FlatMapFunction<String, String>() {  
    public Iterable<String> call(String s) { return Arrays.asList(s.split(" ")); }  
});
```

```
JavaPairRDD<String, Integer> pairs = words.mapToPair(new PairFunction<String, String,  
Integer>() {  
    public Tuple2<String, Integer> call(String s) { return new Tuple2<String, Integer>(s,  
1); }  
});
```

```
JavaPairRDD<String, Integer> counts = pairs.reduceByKey(new Function2<Integer,  
Integer>() {  
    public Integer call(Integer a, Integer b) { return a + b; }  
});
```

```
counts.saveAsTextFile("hdfs://...");
```

Cloud Haskell

Untyped messages

- ★ BSD sockets: low-level streaming API (+ limited datagram oriented API).
- ★ Need higher-level notion of a *message*.
- ★ A message can be arbitrary length.
- ★ For performance and simplicity, multiple segments:
- ★ `type Message = [ByteString]`

network-transport

- ★ A general purpose networking library for *many-process* communication.
- ★ Provides message abstraction.
- ★ Each node has one or more *endpoints*.
- ★ Each process wants *ordered* communication to any other process on any other node.
- ★ Have each process create a connection (or *open a session*) to every other process to encapsulate connection state.
- ★ Under the hood, *multiplex* process-to-process connections into a single transport-level connection.

network-transport (Idealized API)

```
newEndPoint :: Transport -> IO EndPoint
```

```
address      :: EndPoint -> EndPointAddress
```

```
connect      :: EndPointAddress  
              -> IO Connection
```

```
receive      :: EndPoint -> IO Event
```

Processes

```
data Process a
instance Monad Process
instance MonadIO Process

data ProcessId      -- like ThreadId

-- like forkIO
spawnLocal :: Process a -> Process ProcessId
```

Sending/Receiving messages

```
send :: Serializable a  
      => ProcessId  
      -> a  
      -> Process ()
```

```
expect :: Serializable a  
        => Process a
```


Example

```
data Arith = Plus Double Double
           | Mult Double Double
           | Neg Double
```

```
instance Serializable ArithOp
```

```
server :: Process ()
server = forever $ do
    expect >=> say . show . \case of
        Plus x y -> x + y
        Mult x y -> x * y
        Neg x -> negate x
```

```
let'sDoSomeMath :: ProcessId -> Process ()
let'sDoSomeMath there = do
    send there $ Plus 10 2
    send there $ Mult (2^10) (3^10)
    send there $ Neg 1
```

Arith type is ...

- ★ ... indirect;
- ★ ... error prone;
- ★ ... difficult to extend;
- ★ ... antimodular.

Example

```
plus, mult :: Int -> Int -> Process ()  
plus x y = say $ show $ x + y  
mult x y = say $ show $ x * y
```

```
neg :: Int -> Process ()  
neg x = say $ show $ negate x
```

```
let'sDoSomeMath :: NodeId -> Process ()  
let'sDoSomeMath there = do  
    spawnLocal $ plus 10 2  
    spawnLocal $ mult (2^10) (3^10)  
    spawnLocal $ neg 1
```

Remote process

- ★ Message of type `Process` a arbitrary action.
- ★ Not directly serializable (w/o heavyweight runtime support).
 - How to spawn processes remotely (not just locally)?
- ★ **Solution:** track in the type system which actions are serializable!

```
data Closure thing
spawn :: Closure (Process a)
      -> Process ProcessId
```

Applicative forms

$(f \text{ exp}_1 \dots \text{exp}_n)$

In the Identity applicative functor:

```
newtype Id a = Id a
```

```
pure    :: a -> Id a
```

```
(<*>) :: Id (a -> b) -> Id a -> Id b
```

```
pure f <*> pure exp1 <*> ... <*> pure expn
```

(Serializable?) (quasi-)Applicatives

$(f \text{ exp}_1 \dots \text{exp}_n)$

`newtype Closure a = Closure a`

`pure :: a -> Closure a`

`(<*>) :: Closure (a -> b) -> Closure a -> Closure b`

(Serializable?) (quasi-)Applicatives

$(f \text{ exp}_1 \dots \text{exp}_n)$

`newtype Closure a = Closure a`

`pure :: Serializable a => a -> Closure a`

`(<*>) :: Closure (a -> b) -> Closure a -> Closure
b`

Serializable closures

```
cpure :: Serializable a => a -> Closure a  
cap   :: Closure (a -> b)  
      -> Closure a  
      -> Closure b
```

-- A portable code pointer. Like a “name” for expressions.
data StaticPointer a

```
closure :: StaticPointer a -> Closure a
```


-XStaticPointers extension

```
plus, mult :: Int -> Int -> Process ()  
plus x y = say $ show $ x + y  
mult x y = say $ show $ x * y
```

```
plusPtr, mulPtr :: StaticPointer (Int -> Int -> Process ())
```

```
plusPtr = static ptr
```

```
mulPtr = static ptr
```

Rule: can only `static e` if `e` is closed (no free variables).

Example


```
plus, mult :: Int -> Int -> Process ()  
plus x y = say $ show $ x + y  
mult x y = say $ show $ x * y
```

```
neg :: Int -> Process ()  
neg x = say $ show $ negate x
```

```
let'sDoSomeMath :: NodeId -> Process ()  
let'sDoSomeMath there = do  
    spawn there $ closure $ static (plus 10 2)  
    spawn there $ closure $ static (mult (2^10) (3^10))  
    spawn there $ closure $ static (neg 1)
```

Supervision hierarchies

● → "Reports to"



Example

```
plus, mult :: Int -> Int -> Process ()  
plus x y = say $ show $ x + y  
mult x y = say $ show $ x * y
```

```
neg :: Int -> Process ()  
neg x = say $ show $ negate x
```

```
link :: ProcessId  
      -> Process ()
```

```
let'sDoSomeMath :: NodeId -> Process ()  
let'sDoSomeMath there = do  
    pid <- spawn there $ closure $ static (plus 10 2)  
    link pid  
    spawnLink there $ closure $ static (mult (2^10) (3^10))
```

Let's write some code!

Discussion

Discussion

- ★ Distributed programming on the cheap.
 - Captured a useful notion of remote code execution.
 - Minimal runtime support required.
 - How useful is it?
- ★ No good story yet for heterogeneous clusters.
- ★ Cloud Haskell provides low-level mechanism.
 - Build data-parallel models on top
 - Map/Reduce, distributed NDP, DAG of transformers (Spark)

Thank you!